


Metric Learning



16 March 2010

1 Motivation

- 1 Motivation
- 2 Metric Learning for Clustering with Constraints
 - Xing's Method

- 1 Motivation
- 2 Metric Learning for Clustering with Constraints
 - Xing's Method
- 3 Metric Learning for Classification
 - Neighbourhood Components Analysis (NCA)
 - Large Margin Nearest Neighbour (LMNN)

- 1 Motivation
- 2 Metric Learning for Clustering with Constraints
 - Xing's Method
- 3 Metric Learning for Classification
 - Neighbourhood Components Analysis (NCA)
 - Large Margin Nearest Neighbour (LMNN)
- 4 Conclusions

Outline

- 1 Motivation
- 2 Metric Learning for Clustering with Constraints
 - Xing's Method
- 3 Metric Learning for Classification
 - Neighbourhood Components Analysis (NCA)
 - Large Margin Nearest Neighbour (LMNN)
- 4 Conclusions

Introduction

- The notion of (dis-)similarity is fundamental for many data mining / machine learning algorithms, such as:
 - Instance-based supervised algorithms (e.g. kNN)
 - Clustering algorithms (e.g. k-means)
 - Kernel-based algorithms (e.g. SVM)
 - ...
- The traditional approach is to specify the (dis-)similarity a-priori (i.e. before learning)
- This lecture is on adapting (learning) the (dis-)similarities from training data

(Dis-)similarity Learning in Different Settings

Different forms of information to exploit while learning (dis-)similarity:

- Class labels for all the training instances (supervised setting)
- Some constraints on the data (we call it *side-information*)
 - Must-link (ML) and cannot-link (CL) instance level constraints
 - Only ML constraints
 - \mathbf{x}_i is closer to \mathbf{x}_j than to \mathbf{x}_k
 - ...
- No additional information, only the data itself (a.k.a. unsupervised dimensionality reduction or manifold learning)

The Focus in This Lecture

- We will learn Mahalanobis distance metric
- We will exploit information from
 - ML and CL constraints
 - Class labels

The subjects of the next lectures will be:

- Kernel learning
- Unsupervised metric learning
- Feature selection (strongly related with metric learning)

Mahalanobis Metric

- In this work we will learn the Mahalanobis (quadratic) metric parametrized by a symmetric, squared matrix \mathbf{A} :

$$d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i - \mathbf{x}_j)^T \mathbf{A} (\mathbf{x}_i - \mathbf{x}_j)$$

- $\mathbf{x}_i, \mathbf{x}_j \in \mathbb{R}^p$
- $\mathbf{A} \in \mathbb{R}^{p \times p}$
- To keep $d_{\mathbf{A}}$ a valid (pseudo-)metric \mathbf{A} has to be positive demi-definite ($\mathbf{A} \succeq 0$), i.e.
 - $\forall \mathbf{x} \in \mathbb{R}^p \quad \mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$
 - equivalently, \mathbf{A} has non-negative eigenvalues
- Generalizes squared Euclidean metric ($\mathbf{A} = \mathbf{I}$)
- Often \mathbf{A} is the inverse of the data covariance matrix
- \mathbf{A} can be constrained to be diagonal

Mahalanobis Metric \iff Linear Transformation

- We can rewrite $d_{\mathbf{A}}$ using the eigendecomposition of $\mathbf{A} = \mathbf{W}^T \mathbf{W}$ (why is it always possible?):


$$\begin{aligned} d_{\mathbf{W}}^2(\mathbf{x}_i, \mathbf{x}_j) &= (\mathbf{x}_i - \mathbf{x}_j)^T \mathbf{W}^T \mathbf{W} (\mathbf{x}_i - \mathbf{x}_j) \\ &= (\mathbf{W} \mathbf{x}_i - \mathbf{W} \mathbf{x}_j)^T (\mathbf{W} \mathbf{x}_i - \mathbf{W} \mathbf{x}_j) \\ &= (\mathbf{x}'_i - \mathbf{x}'_j)^T (\mathbf{x}'_i - \mathbf{x}'_j) \end{aligned}$$

- $d_{\mathbf{A}}$ is equivalent to applying a simple (spherical) Euclidean metric to the points $\{\mathbf{x}'_i = \mathbf{W} \mathbf{x}_i\}$
- We also start directly with a non-squared $\mathbf{W} \in \mathbb{R}^{d \times p}$, $d < p$


Mahalanobis Metric \iff Linear Transformation (II)

- $d_{\mathbf{A}}$ (and $d_{\mathbf{W}}$) rotates and scales input data
- Good \mathbf{A} (\mathbf{W}) should amplify informative and squash non-informative dimensions
- Possible to define non-linear transformations by exploiting (non-linear) kernel functions - this is not considered in this lecture

Example: Original Data



Example: After Linear Transformation



Outline

- 1 Motivation
- 2 Metric Learning for Clustering with Constraints
 - Xing's Method
- 3 Metric Learning for Classification
 - Neighbourhood Components Analysis (NCA)
 - Large Margin Nearest Neighbour (LMNN)
- 4 Conclusions

Method of Xing et al.'02

Simple idea:

- Keep points in ML as close as possible
- To avoid a trivial solution put some constraint on points in CL

More formally:

$$\min_{\mathbf{A}} \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \text{ML}} d_{\mathbf{A}}^2(\mathbf{x}_i, \mathbf{x}_j) \quad (1)$$

$$\text{s.t.} \quad \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \text{CL}} d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad (2)$$

$$\mathbf{A} \succeq 0 \quad (3)$$

- This optimization problem is *convex*
- Note that (1) is linear while (2) is not
- If we make (2) linear (i.e. use $d_{\mathbf{A}}^2$) then \mathbf{A} will be of rank 1 (similar to LDA)

How to Solve It?

Equivalent formulation

$$\max_{\mathbf{A}} \quad f(\mathbf{A}) = \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \text{CL}} d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j) \quad (4)$$

$$\text{s.t.} \quad g(\mathbf{A}) = \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \text{ML}} d_{\mathbf{A}}^2(\mathbf{x}_i, \mathbf{x}_j) \leq 1 \quad \longrightarrow C_1 \quad (5)$$

$$\mathbf{A} \succeq 0 \quad \longrightarrow C_2 \quad (6)$$


Main idea:

- To optimize (4): gradient ascent
- To satisfy constraints (5) and (6) use the iterative projection algorithm


Iterative Projections

- 1: Initialize \mathbf{A}
 - 2: **repeat**
 - 3: **repeat**
 - 4: $\mathbf{A} \leftarrow \operatorname{argmin}_{\mathbf{A}'} \{ \|\mathbf{A}' - \mathbf{A}\|_F : \mathbf{A}' \in C_1 \}$
 - 5: $\mathbf{A} \leftarrow \operatorname{argmin}_{\mathbf{A}'} \{ \|\mathbf{A}' - \mathbf{A}\|_F : \mathbf{A}' \in C_2 \}$
 - 6: **until** \mathbf{A} converges
 - 7: $\mathbf{A} \leftarrow \mathbf{A} + \nabla_{\mathbf{A}} f(\mathbf{A})$
 - 8: **until** convergence
- First projection (Line 4)
 - (sparse) system of linear equations
 - Second projection (Line 5)
 - $\mathbf{A} = \mathbf{X}^T \mathbf{\Lambda} \mathbf{X}$, $\mathbf{\Lambda} = \operatorname{diag}(\lambda_1, \dots, \lambda_n)$ - matrix of eigenvalues, \mathbf{X} - matrix with eigenvectors as columns
 - $\mathbf{A}' = \mathbf{X}^T \mathbf{\Lambda}' \mathbf{X}$, $\mathbf{\Lambda}' = \operatorname{diag}(\max(\lambda_1, 0), \dots, \max(\lambda_n, 0))$
 - $\|\cdot\|_F$: Frobenious norm over matrices

Iterative Projections: Schematic View



Toy Example



- First plot: original data (three clusters whose centroids differ only in the x and y directions)
- Second plot: Wx for diagonal A
- Third plot: Wx for full A

Example and Figure taken from Xing et al. '02

Xing's method - conclusions

Pros:

- Simple idea
- Improves performance of standard k-means
- Can be used in supervised setting

Cons:

- Complex to solve
- Iterative projections can be computationally expensive

Outline

- 1 Motivation
- 2 Metric Learning for Clustering with Constraints
 - Xing's Method
- 3 Metric Learning for Classification
 - Neighbourhood Components Analysis (NCA)
 - Large Margin Nearest Neighbour (LMNN)
- 4 Conclusions

kNN

We focus on kNN since:

- It is easy to implement
- Decision boundaries are non-linear
- Its complexity does not depend on the number of classes
- In the distance metric is carefully selected : state-of-the-art results
- Only one parameter to tune (k)
- Some theoretical guaranties of good performance

Neighbourhood Components Analysis (NCA)

- We want to improve the performance of kNN by learning metric
 - the metric should reduce generalization error (E_G)
- We will approximate the generalization error by the error estimated using leave-one-out cross validation procedure (E_{LOO})
 - We assume $E_{LOO} \approx E_G$
 - The goal is then to keep E_{LOO} as small possible
- However, E_{LOO} is a difficult objective function to optimize with respect to the metric
 - *highly discontinuous* function of the metric

NCA - Stochastic Neighbourhood Selection

Main idea:

- 1 For each instance select the neighbours in a *stochastic manner*
- 2 Look for the *expected votes* for each class

More formally:

- 1 Each instance \mathbf{x}_i select another instance \mathbf{x}_j as its neighbour with probability p_{ij}
 - p_{ij} should be higher for points that are closer according to $d_{\mathbf{W}}(\mathbf{x}_i, \mathbf{x}_j)$

$$p_{ij} = \frac{\exp -d_{\mathbf{W}}^2(\mathbf{x}_i, \mathbf{x}_j)}{\sum_{k \neq i} \exp -d_{\mathbf{W}}^2(\mathbf{x}_i, \mathbf{x}_k)}, p_{ii} = 0$$

- p_{ij} quickly reaches 0 expect for points that are close
- 2 The probability p_i that instance \mathbf{x}_i will be correctly classified:

$$p_i = \sum_{\{j : y_i = y_j\}} p_{ij}$$

NCA - Cost Function

- The expected (soft) leave-one-out classification performance:


$$\begin{aligned}
 f(\mathbf{W}) &= \frac{1}{n} \sum_i p_i \\
 &= \frac{1}{n} \sum_i \sum_{\{j : y_i=y_j\}} p_{ij} \\
 &= \frac{1}{n} \sum_i \sum_{\{j : y_i=y_j\}} \frac{\exp -d_{\mathbf{W}}^2(\mathbf{x}_i, \mathbf{x}_j)}{\sum_{k \neq i} \exp -d_{\mathbf{W}}^2(\mathbf{x}_i, \mathbf{x}_k)}
 \end{aligned}$$

- $f(\mathbf{W})$ is continuous and differentiable, but not convex
- The goal is to find \mathbf{W} that maximizes $f(\mathbf{W})$, i.e.

$$\max_{\mathbf{W}} f(\mathbf{W})$$

- We *do not need* the $\mathbf{W} \succeq 0$ constraint


Toy Example



- Original data (in \mathbb{R}^3): Five classes organized in concentric rings in \mathbb{R}^2 , the third dimension is an uncorrelated noise
- First plot: PCA
- Second plot: LDA
- Third plot: NCA (we learn $\mathbf{W} \in \mathbb{R}^{2 \times 3}$)

Example and Figure taken from Glodberger et al. '05

Real World Example



- Original data (in \mathbb{R}^{560}): images of 18 faces (100 images per face)
- First plot: PCA
- Second plot: LDA
- Third plot: NCA (we learn $\mathbf{W} \in \mathbb{R}^{2 \times 560}$)

Example and Figure taken from Glodberger et al. '05

NCA - Conclusions

Pros

- Usually good performance
- No assumptions about the data such that unimodality, linear separability, etc.

Cons:

- Not convex
- Computational complexity
 - $O(n^2p^2)$ for full matrices \mathbf{W}
 - $O(n^2p)$ for rectangular "thin" matrices \mathbf{W}

Large Margin Nearest Neighbour (LMNN)

- For each learning instance set a number of closest instances (*target neighbours*) that share the same class
- The target neighbours are defined using standard Euclidean metric
- Similarly to NCA optimizes a surrogate function for LOO error

For each instance:

- 1 Pull target neighbours closer
- 2 Push away instances of different classes that are in the "neighbourhood" (i.e. *impostors*)
 - By a large margin

LMNN - Schematic View






Figure taken from Weinberger et al. '06

Notation

- $\eta_{ij} = \begin{cases} 1 & \text{if } \mathbf{x}_i \text{ is a target neighbour of } \mathbf{x}_j, \\ 0 & \text{otherwise} \end{cases}$
 - need to be specified before learning
- $y_{ij} = \begin{cases} 1 & \text{if } y_i = y_j, \\ 0 & \text{otherwise} \end{cases}$
- $[k]_+ = \max(k, 0)$: hinge loss




Formal Definition



$$\min_{\mathbf{A}} \{ \text{pull}(\mathbf{A}) + \text{push}(\mathbf{A}) \}$$

- $\text{pull}(\mathbf{A}) = \sum_{ij} \eta_{ij} d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j)$
 - pull together the neighbours
- $\text{push}(\mathbf{A}) = \sum_{ijl} \eta_{ij}(1 - y_{il}) [1 + d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j) - d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_l)]_+$
 - Push away impostors by a margin of 1
 - hinge loss



Formal Definition (2)

$$\begin{aligned}
 \min_{\mathbf{A}} \quad & \sum_{ij} \eta_{ij} d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j) \\
 & + \sum_{ijl} \eta_{ij} (1 - y_{il}) [1 + d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j) - d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_l)]_+ \\
 \text{s.t.} \quad & \mathbf{A} \succeq 0
 \end{aligned}$$

Equivalent semi-definite program (convex):

$$\begin{aligned}
 \min_{\mathbf{A}} \quad & \sum_{ij} \eta_{ij} d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j) + \sum_{ijl} \eta_{ij} (1 - y_{il}) \xi_{ijl} \\
 \text{s.t.} \quad & d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_j) - d_{\mathbf{A}}(\mathbf{x}_i, \mathbf{x}_l) \geq 1 - \xi_{ijl} \\
 & \xi_{ijl} \geq 0, \mathbf{A} \succeq 0
 \end{aligned}$$

Possible to solve it efficiently even though the number of constraints scales as $O(n^2k)$

- Possible to solve problems with $n \approx 60,000$ in a reasonable time

LMNN - Some Empirical Results

Estimated error:

	MNIST	News	Isolet	Bal	Faces	Wine	Iris
p	784	30000	617	4	1178	13	4
p after PCA	164	200	172	-	30	-	-
n (train)	60000	16000	6238	445	280	126	106
Euclidean kNN	1.9	20.0	9.4	14.1	8.2	30.0	4.3
LMNN kNN	1.2	11.0	4.7	10.0	0.3	1.1	3.5

Results taken from Weinberger et al. '06

LMNN - Conclusions

Pros

- State-of-the-art metric learning results
- No assumptions about the data such that unimodality, linear separability, etc.
- Convex problem
- Can be solved very efficiently
- Because of the hinge-loss it is similar to SVM

Cons:

- Sensitive to outliers
- Target neighbours need to be specified a-priori
- Too many constraints - this might affect the quality of the solution

Outline

- 1 Motivation
- 2 Metric Learning for Clustering with Constraints
 - Xing's Method
- 3 Metric Learning for Classification
 - Neighbourhood Components Analysis (NCA)
 - Large Margin Nearest Neighbour (LMNN)
- 4 Conclusions

Metric Learning - Conclusions

Pros

- Unified framework for learning a linear transformation that better separates classes
- In most cases improves over standard Euclidean metric
- Often reported state-of-the-art results
- Can be non-linear (not discussed here)
- Can be very efficient

Cons:

- Scales as $O(p^2)$ ($O(p)$) - so data usually preprocessed by e.g. PCA
- Convex methods are less flexible; flexible methods and non-convex